


(a) Router-level alignment intervention  Harmful Experts  Safe Experts  Neutral Experts



(b) Different risky categories engage different expert subsets



(c) Our method escapes the safety-helpfulness trade-off

